




Variables Tested

- ✓ **Verification affordance**
(governs override)
- ~ Task complexity
- ~ Error detectability
(continuous modulation)
- ↓ **Model capability**
(modulates, not eliminates)
- ✗ Trace quality
- ✗ Source authority
- ✗ Model confidence
(null predictors)

On verification-costly tasks, CIF reaches 61% with 100% binary switching

On verifiable tasks, CIF $\approx$ 0% — traces are harmless or beneficial